Friday: Project Notes

Friday is a voice-first personal assistant built by Aditya in Bengaluru. It runs on his own laptop and is controlled by holding the Right Ctrl key while speaking, which is called push-to-talk.

For speech to text, Friday uses faster-whisper with the base.en model. For text to speech, it uses Piper with the en_US-amy-medium voice. Piper speaks the reply one sentence at a time.

The main language model is Groq running gpt-oss-20b. If Groq fails, Friday falls back to Gemini. Both are called through one file named llm.py.

Friday uses LangGraph to run each command. The graph goes from the router to a safety check, then to either a skill or the language model, and finally to the reply.

All data is stored in SQLite. This includes chat history, reminders, lists, traces and documents.

Reminders are scheduled with APScheduler and saved in SQLite, so they survive a restart. Friday cannot yet repeat reminders, list them or cancel them.

The music skill uses the Spotify Web API. It plays only on the device named ADI7 and never falls back to another device. Friday cannot yet play playlists or albums by name, and it cannot set the volume by spoken numbers like volume fifty.

The chat skill remembers the last six turns of conversation. A turn is one message from the user plus Friday’s reply.

Math is handled by a safe parser that never uses eval. Lists and notes are stored in a table called list_items.

The catch skill is always registered last. If a command sounds like a list, reminder or timer command but is garbled, it asks the user to say it again instead of letting the language model pretend it saved something.

Secrets such as API keys live only in the .env file and the .secrets folder, and both are ignored by git.